





# MINECRAFT FPGA EDITION



**6.2050 Project Proposal Presentation**  
Yohan Guyomard, Win Win Tjong  
November 5, 2024


# Project Overview

- ❖ FPGA client that connects to Minecraft worlds
  - Interact with the same world as other players\*
- ❖ Renderer uses a ray-tracing approach, ideally with several bounces for realistic effects like shadows, reflections and refraction.



\*\*custom plug-in is used to replace Minecraft's protocol (and TCP) with UART

# Terminology & Preliminary Stuff




voxel (or block)



chunk


Segment of world stored  
on the FPGA ( $128^3$  voxels)



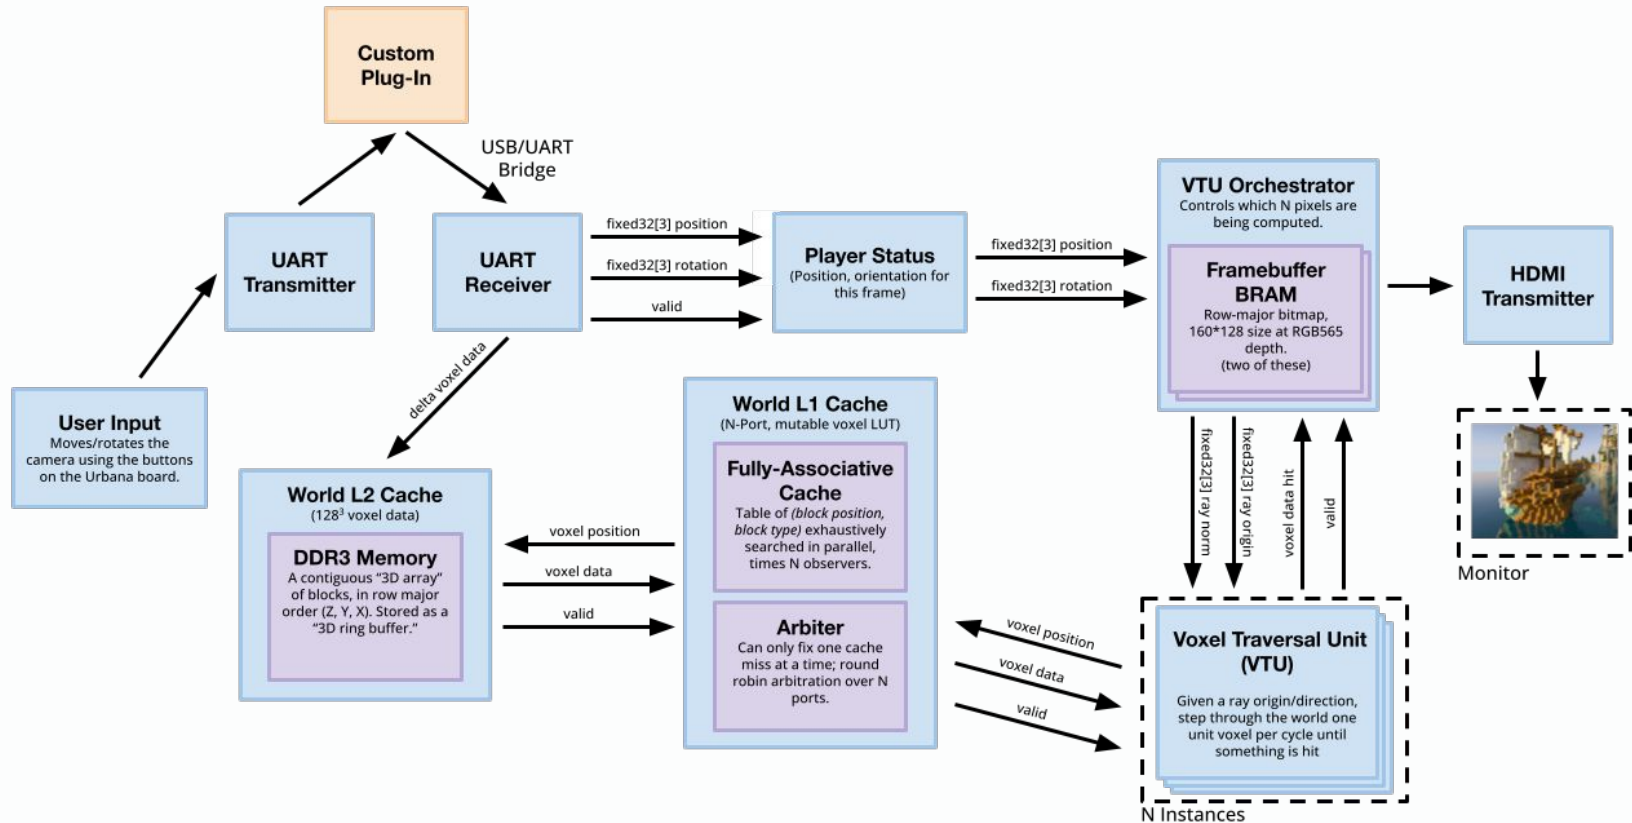
voxel traversal

Steps one voxel at a time (air=continue; block=halt)  
(Amanatides & Woo)

Ray-tracing




# Block Diagram



# L2 Cache

- ❖ Choice of data structure to store voxel data.
  - Update happens when player status changes
    - Direction depends on whichever direction the player moves towards
  - Stored in DDR3
  - Each voxel is represented with 5 bits
  - Total size is  $128^3 * 5 \approx 10\text{Mb}$





# UART

- ❖ Communication protocol between the game's server to the FPGA.
  - The asynchronization between transmission and receiving allows for flexibility when the two are happening at different rates.
  - Straightforward to use and allows for easier debugging.
  - Receives data from the game plug-in and stores it directly into the L2 cache.
    - The plug-in used for this project should be a simple implementation and should not hinder other parts of the project.
  - To run a smooth 30 fps, rate of transfer exceeds  $128^2 \cdot 5 \cdot 30 \approx 2 \text{ Mb}$ .
    - The UART module in the Urbana board supports a transfer rate up to 2Mb, so it should be fine.

# Player Input & Status

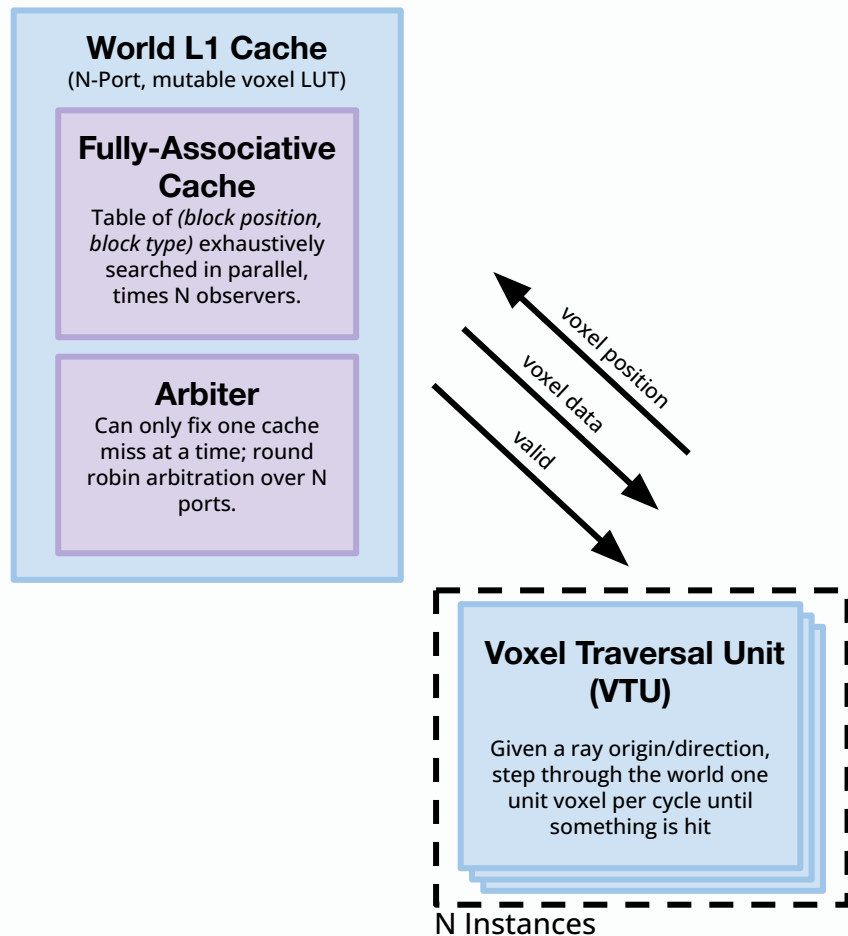
- ❖ The input of a joystick connected to the board are sent to the game via UART transmitter, which will update player status in-game.
  - This then is transmitted into the UART receiver and into the L2 cache
  - At the same time, this will update player's position, which will be accessed by the frame buffer.



© Photo by ElectroPunk

# Memory Bottleneck


- We can compute the path of the ray really fast and in-parallel
- **But** *each* iteration of *each* VTU needs 1 memory access
  - Voxel is air? continue.
  - Voxel is not? halt.
- L2 cache has just **one** reading port, so really it's just 1 working VTU with deadtime for  $N - 1$  VTUs





# World L1 Cache

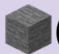
- ❖ Likely, all rays will query the same set of voxels\*
- ❖ This data structure has **N ports**, each simultaneous single-cycle access if cache-hit!
- ❖ On cache miss, it's potentially answering many VTUs' requests all at once
  - If not, arbiter takes over
  - Continues to serve the non-blocked ports!



\*educated speculation

# World L1 Cache (Problem)

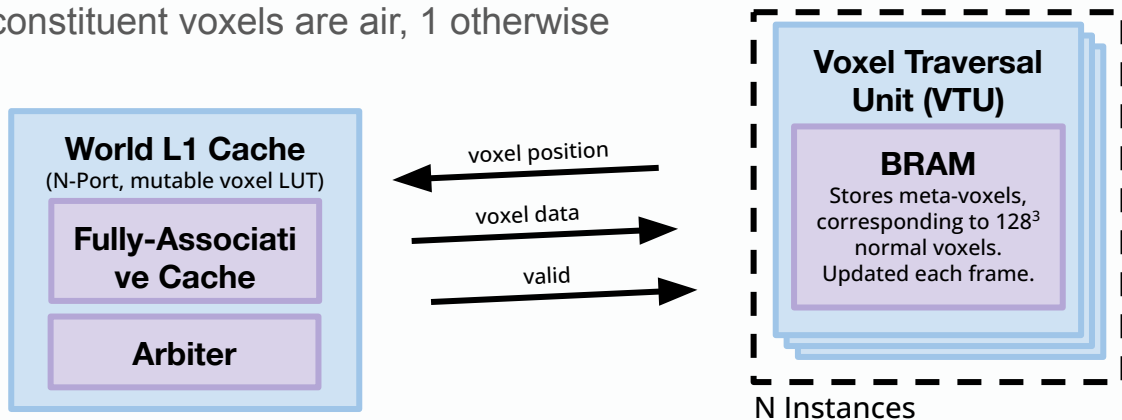
- ❖ Can't make this very large\*
  - For N ports and M entries, that's  $N \cdot M$  comparators
  - There's a reason K-way set associative caches exist, but these don't really apply here
- ❖ The bulk of the cache will just be "air" voxels\*
  - Wasteful; "air" or "not air" is 1 bit of information, not an entire voxel ID!

L1 Cache	
(0, 1, 5)	(air)
(-3, 0, 0)	 (1)
(4, 1, -2)	(air)
(0, 0, 0)	(air)

\*educated speculation

# Sparseness Cache

- ❖ Each VTU has its own cache of the entire chunk!
  - 1 BRAM per VTU
- ❖ Only care about air/not air, so  $128^3$  bits of data
  - Actually that doesn't fit...  $\sim 2$  Mbits  $>$  32 Kbits
- ❖ So, create “meta-voxels” and store that!
  - 1 meta-voxel =  $4^3$  voxels
  - Occupies 1 bit: 0 if all its constituent voxels are air, 1 otherwise



# VTU Orchestrator

- ❖ One more memory thing: how do N VTUs draw to the same framebuffer?
  - The orchestrator will block a VTU until its pixel can be pushed to the framebuffer BRAM
  - Shouldn't create much deadtime, the VTU are all asynchronous anyways

# Project Checkoff List

## ❖ The Commitment

- Implement a basic renderer using ray-tracing
- Stream a static world from a Minecraft server onto the FPGA


## ❖ The Goal

- Sufficiently optimize our renderer for higher framerates, (more) light bounces, textures; overall something more pleasing to the eyes.
- Continuously update the FPGAs copy of the voxel world with incoming data; from a user's point of view, the world should seem infinite.

## ❖ The Stretch

- Higher resolution graphics (going beyond BRAM) and/or higher framerates
- Rendering entities (players, mobs, etc.)

# Sparseness Cache



# Timeline

	<b>Computer Interface</b> (Win Win)	<b>Rendering</b> (Yohan)
<b>Week 0</b> (Nov. 6)	Implement L2 cache (ring)	Implement L1 cache
<b>Week 1</b> (Nov. 11)	Integrate into L1 cache	Implement VTU
<b>Week 2</b> (Nov. 18)	UART transceiver (protocol)	Add occupancy optimization
<b>Week 3</b> (Nov. 25)	User input and player status	Create VTU orchestrator
<b>Week 4</b> (Dec. 2)	Final integration	Final integration
<b>Week 5</b> (Dec. 9)	Finalize report	Finalize report

**Thank you! Questions?**



# Update: Plugin has been implemented

- ❖ We are able to query blocks in a Minecraft world
- ❖ From the same program, we have serial communication (currently tested with a microcontroller)



← Joe, right now

# whaterweworkinwit

		I/O Optimization at the Lowest Cost and Highest Performance-per-Watt (1.0V, 0.95V)					
	Device Name	XC7S6	XC7S15	XC7S25	XC7S50	XC7S75	XC7S100
Logic Resources	Logic Cells	6,000	12,800	23,360	52,160	76,800	102,400
	Slices	938	2,000	3,650	8,150	12,000	16,000
	CLB Flip-Flops	7,500	16,000	29,200	65,200	96,000	128,000
Memory Resources	Max. Distributed RAM (Kb)	70	150	313	600	832	1,100
	Block RAM/FIFO w/ ECC (36 Kb each)	5	10	45	75	90	120
	Total Block RAM (Kb)	180	360	1,620	2,700	3,240	4,320
Clock Resources	Clock Mgmt Tiles (1 MMCM + 1 PLL)	2	2	3	5	8	8
I/O Resources	Max. Single-Ended I/O Pins	100	100	150	250	400	400
	Max. Differential I/O Pairs	48	48	72	120	192	192
Embedded Hard IP Resources	DSP Slices	10	20	80	120	140	160
	Analog Mixed Signal (AMS) / XADC	0	0	1	1	1	1
	Configuration AES / HMAC Blocks	0	0	1	1	1	1
Speed Grades	Commercial Temp (C)	-1,-2	-1,-2	-1,-2	-1,-2	-1,-2	-1,-2
	Industrial Temp (I)	-1,-2,-1L	-1,-2,-1L	-1,-2,-1L	-1,-2,-1L	-1,-2,-1L	-1,-2,-1L
	Expanded Temp (Q)	-1	-1	-1	-1	-1	-1

- based on feedback from the block diagram and some results from our initial implementation, a large-enough, single-cycle useful L1 cache may be unrealistic.
- What if each VTU had its own copy of the voxel world?
  - Can't each have their own DDR RAM
  - Can't fit all the data we need in BRAM
  - But, for 99% of memory accesses we only care about "voxel exists or is it air?"
    - Reduce to 1-bit representation of the world
  - $128^3 = 2$  Mbits is still too large
  - But, sample every four voxel (i.e. 1 if any of those  $4^3$  voxels are not air, 0 otherwise), one BRAM is *just* enough (32 Kbits)
- There are other, more compact ways to represent sparseness (i.e. octree) but less trivial to implement
  - Maybe a stretch goal

***basically, less memory accesses = less contention = less bad***

## can we rethink the L1 cache?

- now that each VTU has its own “approximate” copy of the world, is it better to change the L1 cache to be e.g. BRAM with more entries but slower access?
- We would need an arbiter, one VTU reads L1 cache at a time

*basically, less memory accesses = less contention = less bad*

→ **can we race the beam, is it possible?**

- Without cheating, definitely not. Hell no. 720p @ 60FPS means 74.25 MHz clock and we can't trace an entire ray in one clock cycle
- *But*, our ray tracing algorithm is upper bounded
  - $\text{max}(\# \text{ cycles}) = (\# \text{ bounces}) * (\text{max travel distance})$
  - $(\text{max travel distance}) = 221$  for  $128^3$  chunk
- so when we're displaying pixel  $i$ , we need to start rendering pixel  $i + 221$ , and  $i + 220$ ,  $i + 219$ , etc. had started in prior cycles.
- so we need ~220 VTUs... not gonna happen

***No, pretty much not possible***

→ **with a framebuffer, how large can it get?**

- (assuming storage, to/fro is a non-issue)
- 221 max ray distance (VTU steps one voxel at a time)
- assuming single-cycle step (not realistic):
  - $100\text{MHz} / 221 = \sim 452,000$  pixels per VTU per second
  - Comfortably draw 480p@30 with 20 VTUs, 720p with 60 VTUs, etc.
- more realistically, two-cycle step average over  $221/4 = 55$  meta-voxels
  - Thrown in there are “meta-voxel misses” (thinks it hit solid, but had to go check that it was air) which can cost a lot of cycles if it’s an L1 cache-miss. No idea how to estimate this, let’s say it quadruples the worst case average.
  - $100\text{MHz} / (55 * 4) = \sim 454,000$  pixels per VTU per second
- if we really limit the size of the frame buffer (e.g. a cute  $128*160$ ), we can *really* go crazy with light bounces
  - At this size, just one VTU is enough for everything (single-ray)

***720p is possible, but choosing lower resolutions gives more room for creativity & artistic effects!!***